Your resume.
Your next step.

Analysis date (UTC): 2026-09-28T18:36:46+00:00

Document: data_analyst.pdf

Target: Machine Learning Engineer

25.5% Skill alignment	8 Skills detected	4 Core skills not detected

Match scores are estimates of documented skill alignment, not recruiter decisions or hiring probabilities. This tool supports career guidance, never automatic hiring or rejection. Not detected in the resume does not mean a skill is absent.

Recommended roles

1. Data Analyst - 95.5% (5/5 core skills detected)
2. Business Intelligence Analyst - 70.5% (4/5 core skills detected)
3. Data Scientist - 59.8% (4/6 core skills detected)

Target-role skill gaps

Core detected: Python

Core not detected in the resume: Machine Learning, scikit-learn, FastAPI, Docker

Preferred detected: Pandas, SQL

Preferred not detected: Model Evaluation, MLflow, Cloud Deployment

All detected skills

BI Tools: Power BI

Data Analysis: Data Cleaning, Data Visualization, Excel, Pandas, Statistics

Databases: SQL

Programming Languages: Python

Evidence strength

Project/work-section mentions: Power BI, Data Cleaning, Data Visualization, Excel, Pandas, Statistics, SQL, Python

Mention only; add evidence where truthful: None

These labels describe where terms appear, not independently verified skill levels. Repetition does not increase the score.

Personalized learning plan

Each week is a suggested study block of roughly 5-8 hours. Adjust the pace to your background. Dependencies may add foundation topics; this is not a promise of mastery.

Week 1: Machine Learning

Core requirement. Train/test splits, baselines and model selection

Practice: Compare a baseline and a classifier without test-set leakage.

Official learning resource

Week 2: scikit-learn

Core requirement. Pipelines, preprocessing and cross-validation

Practice: Build a preprocessing-plus-model pipeline with held-out evaluation.

Official learning resource

Week 3: REST APIs

Foundation for a later topic. HTTP methods, status codes and resource design

Practice: Design a CRUD API and document requests and error responses.

Official learning resource

Week 4: FastAPI

Core requirement. Typed request schemas and API documentation

Practice: Serve a model through a validated endpoint and test it.

Official learning resource

Week 5: Linux

Foundation for a later topic. Files, processes, permissions and shell basics

Practice: Run an application as an unprivileged user and inspect its logs.

Official learning resource

Week 6: Docker

Core requirement. Images, containers, volumes and non-root execution

Practice: Containerize a service with a health check and no embedded secrets.

Official learning resource

Week 7: Applied project

Put the target role into practice. Build a portfolio artifact for Machine Learning Engineer.

Practice: Train a reproducible classifier with a held-out evaluation, serve predictions through FastAPI, and containerize the tested service.

Resume improvements

Use a measurable outcome when available

No common outcome metric was found in the detected project/experience sections. Add truthful measures such as dataset size, accuracy with validation method, or time saved; do not invent numbers.

Link to work you can demonstrate

Consider a portfolio or repository link for a relevant project. Do not include confidential code or data.

Check the target-role gaps

Not detected: Machine Learning, scikit-learn, FastAPI, Docker. If you already use these skills, add truthful examples; otherwise follow the learning plan.

How your score is calculated

With preferred skills: 75 x core coverage + 15 x preferred coverage + 10 x TF-IDF cosine similarity. Without preferred skills: 90 x core coverage + 10 x similarity. Coverage and similarity are fractions from 0 to 1.

Core: 15.00 points; preferred: 6.00 points; similarity: 4.54 points. Total rounded to one decimal: 25.5.

Both TF-IDF vectors contain canonical, role-relevant skill tokens only. The resume vector is projected onto the target's skills. The fixed vocabulary and IDF values are learned from the 13 curated role profiles, not from personal resume text. This is lexical skill similarity, not semantic understanding.

Gender, age, religion, nationality, photographs, marital status and disability are not score features. Recognized personal fields and sections are excluded from skill extraction. Education and section context inform the displayed analysis but do not change the score. The dictionary and role templates are educational, not employer-validated hiring criteria.

Privacy

This report contains derived analysis only, not the full resume text. Uploaded bytes and results are held in the active session's memory and are not written to an application database. Use Clear session to discard them. Your browser download is your own copy.